

GLS versus Averaged Change-Score Estimators for Clinical Trials with Run-In Observations

Ronald G. Thomas, Ph.D.*

August 15, 2026

Abstract

Background. Three estimators of the treatment effect in longitudinal trials with run-in observations are in routine use: the GLS estimator under the linear mixed model, which exploits the full covariance structure to extract maximum information from all observations; the ANCOVA estimator, which regresses post-randomization means on the run-in mean as a covariate without specifying the random-effects covariance; and the averaged change-score estimator, which computes a simple difference of phase means and applies a t -test, requiring no variance-component estimation. The relative efficiency of these estimators under correct and misspecified models has not been comprehensively characterized. **Methods.** We derive the variance of the three estimators as a function of the trial design (J_0, J_1, J_2) and the variance parameters (σ^2, σ_b^2) . The derivation proceeds via the Woodbury-based variance formulas of Paper 1 in this compendium, plus closed-form expressions for the ANCOVA and averaged-change-score variances under the same random-slopes model. Robustness to misspecification of the within-subject covariance (R) and the random-effects covariance (G) is assessed analytically through sandwich variances, and a Monte Carlo simulation under compound symmetry evaluates the estimators alongside a GLS analysis that misstates the random-slope variance. **Results.** All comparisons are made on the common estimand scale, with the averaged and ANCOVA contrasts rescaled to unbiasedness for the treatment effect γ . Under correct specification, $\text{GLS} \leq \text{ANCOVA} \leq \text{averaged}$ in variance, with no violations over a 115-cell grid of designs and signal-to-noise ratios. The efficiency penalty of the simpler estimators is substantial in run-in designs with multiple treatment-period visits: at $(J_0, J_1, J_2) = (2, 2, 0)$ the relative efficiency of ANCOVA versus GLS is 0.64 to 0.81 across signal-to-noise ratios, an efficiency loss of 19% to 36%. Relative efficiency above 0.95 occurs only in single-post-visit run-in designs ($J_1 = 1, J_2 = 0$) and in designs without run-in or common close ($J_0 = 0, J_2 = 0$) at high signal-to-noise. In the MIRIAD-based example the ANCOVA sample size exceeds the GLS sample size by up to a factor of roughly 2.6 among run-in designs, and the averaged sample size by up to a factor of roughly 45. Under the G and R misspecifications studied, GLS never loses its advantage over the simpler estimators: the misspecified-GLS sandwich variance peaks at 5.3 against an averaged-estimator variance of 8.9 at the reference design. **Conclusions.** GLS is the efficient estimator under correct specification, and its advantage survives the covariance misspecifications examined here. The averaged change-score estimator remains useful as a fully prespecified robust benchmark and a teaching tool. The open robustness question, not studied here, is whether feasible GLS with variance components estimated from small samples retains this advantage.

*Wertheim School of Public Health, UCSD; ORCID: 0000-0003-1686-4965

40	Contents	
41	1 Introduction	2
42	2 Three estimators	4
43	2.1 GLS under the mixed model	4
44	2.2 Averaged change score (<i>t</i> -test)	4
45	2.3 ANCOVA with run-in mean as covariate	5
46	2.4 SUR (seemingly unrelated regression)	6
47	2.5 BLUP as theoretical benchmark	6
48	3 Relative efficiency under correct specification	7
49	3.1 Analytical comparison	7
50	3.2 When does the gap exceed 5%?	8
51	4 Robustness to misspecification	10
52	4.1 Misspecified <i>G</i>	10
53	4.2 Misspecified <i>R</i>	10
54	5 Simulation study	10
55	5.1 Design	10
56	5.2 Results	12
57	6 Numerical examples	13
58	6.1 MIRIAD parameters	13
59	6.2 When to use which estimator	13
60	7 Discussion	13
61	8 Conflict of interest	14
62	9 Funding	15
63	10 Data availability statement	15
64	11 Reproducibility	15
65	12 References	15

66 1 Introduction

67 The GLS estimator $\hat{\gamma}_{\text{GLS}} = (X'\Sigma^{-1}X)^{-1}X'\Sigma^{-1}\mathbf{c}$ is the best linear unbiased esti-
68 mator (minimum variance among linear unbiased estimators) when Σ is known.
69 In practice, Σ must be estimated, and the efficiency of the feasible GLS depends
70 on the accuracy of this estimation. With small sample sizes or short follow-up,
71 variance component estimates can be unstable, potentially eroding the theoretical
72 advantage of GLS.

73 The averaged change-score approach sidesteps variance component estimation entirely. For a

trial with J_0 run-in and J_2 common close observations, the estimator is:

$$\hat{\gamma}_{\text{avg}} = \bar{d}_1 - \bar{d}_0 \quad (1)$$

where

$$d_{ik} = \frac{1}{J_1 + J_2} \sum_{\ell=1}^{J_1+J_2} Y_{i,\ell,k} - \frac{1}{J_0} \sum_{\ell=1}^{J_0} Y_{i,-\ell,k} \quad (2)$$

is the difference of the post-randomization and run-in means for participant i . The baseline visit is the change-score reference and cancels, so it enters neither average; this matches the contrast implemented in `var_gamma_avg()`. Throughout, the code and figure parameters use (r, k, f) for the design counts, with the correspondence $J_0 = r$ (run-in), $J_1 = k$ (treatment period), and $J_2 = f$ (common close); the prose and table headers use (J_0, J_1, J_2) .

The debate between change-score and ANCOVA analyses has a long history^{1-3,4} provided a formal efficiency comparison for pretest-posttest designs. The closest antecedent of the present work is⁵, who derived the relative efficiency of post-treatment, change-score, and ANCOVA analyses when several baseline measurements are available, and⁶, who extended the treatment to repeated post-treatment measurements. Their ordering, that ANCOVA on the mean of the available baselines dominates the change score which in turn dominates the raw post-treatment comparison, is the qualitative result that the ordering established in Section 3.1 specializes. What is new here is not the ordering itself but the closed-form relative efficiencies in the run-in random-slope parameterization: the efficiency margins are obtained as explicit functions of the design counts (J_0, J_1, J_2) and the signal-to-noise ratio σ_b/σ , they cover a post-treatment common close phase absent from the Frison-Pocock setting, and the estimand is a difference of slopes rather than a difference of means. Within the run-in literature,⁷ derived GLS power formulas for trials with pre-randomization observations, which the present comparison takes as the efficient benchmark against which the simpler estimators are measured.⁸ and⁹ compared baseline-as-covariate against baseline-as-response (constrained longitudinal data analysis).¹⁰ developed a semiparametric framework for covariate adjustment.¹¹ showed that MMRM requires time-by-covariate interactions to guarantee power gains over ANCOVA.¹² established the foundational comparison of pretest-posttest designs.¹³ addressed conditioning on baseline in longitudinal RCTs.¹⁴ proposed the two-period approach for run-in trials.

We note that the averaged change-score estimator studied here is distinct from (and less efficient than) the constrained longitudinal data analysis (cLDA) estimator of¹⁵ and⁹, which includes baseline in the response vector with a common-mean constraint. The cLDA estimator uses the full covariance structure without requiring specification of the random effects distribution G , placing it between the averaged estimator and GLS in the efficiency ordering.¹⁶ extended the cLDA framework to pre- and post-treatment measurements with equal baseline assumptions. A full comparison including cLDA is beyond the scope of this paper but represents

113 a natural extension.

114 The `longpower` R package¹⁷ provides the CRAN reference implementation of
115 power calculations for the GLS / REML-based slope-comparison estimator, cov-
116 ering the Liu-Liang¹⁸, Diggle¹⁹, and Edland²⁰ formulations. These routines corre-
117 spond to the $\hat{\gamma}_{\text{GLS}}$ estimator studied here under correct specification and constitute
118 the Paper 1 baseline extended by run-in and common close observations. However,
119 `longpower` does not cover the ANCOVA-on-run-in-mean estimator or the averaged
120 change-score t -test, nor does it characterize the efficiency loss from misspecifica-
121 tion of G or R . The comparison developed here therefore complements `longpower`
122 by (i) ordering the three candidate estimators under correct specification and (ii)
123 quantifying the efficiency margin that GLS retains under the misspecifications of
124 G and R studied below.

125 2 Three estimators

126 2.1 GLS under the mixed model

127 From²¹, the per-pair variance is the (γ, γ) element of $(X'\Sigma^{-1}X)^{-1}$ with $\Sigma =$
128 $R + ZGZ'^{22,23}$, computed via the Woodbury identity.

129 For the Frost conventional design ($J_0 = 0$, $J_1 = 2$):

$$130 \quad \hat{\gamma}_{\text{GLS}} = \frac{1}{2t}[(C_{22} + C_{21}) - (C_{12} + C_{11})] \quad (3)$$

131 with $\text{Var}(\hat{\gamma}_{\text{GLS}}) = (\sigma^2 + 2\sigma_b^2 t^2)/t^2$.

132 2.2 Averaged change score (t -test)

133 The averaged change-score estimator ignores the within-subject correlation structure and
134 treats d_{ik} as independent observations:

$$135 \quad \text{Var}(\hat{\gamma}_{\text{avg}}) = \frac{2}{m} \text{Var}(d_{ik}) \quad (4)$$

136 The variance $\text{Var}(d_{ik})$ depends on the correlation structure but does not require its estimation.
137 As an estimator of γ , the contrast is divided by the design constant c derived in Section 3.1,
138 and all variances, relative efficiencies, and sample sizes reported in this paper are on that
139 common γ scale.

```
a <- 1.0
b <- 1.0
tt <- 1.0

# gamma-scale rescaling constant (Section 3.1):
# mean accumulated effect over post-randomization visits
```

```

cfac <- function(r, k, f, t_interval = 1) {
  t_interval * (k * (k + 1) / 2 + k * f) / (k + f)
}

v_gls <- var_gamma_matrix(2, 2, 0, a, b, tt)
v_avg <- var_gamma_avg(2, 2, 0, a, b, tt) /
  cfac(2, 2, 0, tt)^2
re <- v_gls / v_avg

cat(sprintf('J_0=2, J_1=2: GLS=%.4f, Avg=%.4f, RE=%.4f\n',
  v_gls, v_avg, re))

```

```
140 ## J_0=2, J_1=2: GLS=2.1569, Avg=8.8889, RE=0.2426
```

141 2.3 ANCOVA with run-in mean as covariate

142 The ANCOVA approach uses the averaged run-in rate as a baseline covariate without specifying σ_b^2 :

$$144 \bar{Y}_{i,\text{post}} = \alpha + \gamma g_i + \phi \bar{Y}_{i,\text{pre}} + \varepsilon_i \quad (5)$$

145 The variance of $\hat{\gamma}$ under this model is

$$146 \text{Var}(\hat{\gamma}_{\text{ANC}}) = \frac{2}{m} \text{Var}(\bar{Y}_{\text{post}}) (1 - \rho_{\text{pre,post}}^2) \quad (6)$$

147 where $\rho_{\text{pre,post}}$ is the marginal correlation between the run-in and post-randomization means. When $J_0 = 0$ there is no covariate and ANCOVA reduces to the averaged estimator. When $\rho_{\text{pre,post}}$ is large (i.e., when σ_b^2 dominates σ^2), the $(1 - \rho^2)$ factor produces substantial variance reduction.

151 The key advantage over GLS is that ϕ is estimated from the data rather than derived from assumed variance components^{1,24}. The key advantage over the averaged estimator is that it exploits the predictive value of the run-in observations^{2,25}.

```

c220 <- cfac(2, 2, 0, tt)
v_gls <- var_gamma_matrix(2, 2, 0, a, b, tt)
v_anc <- var_gamma_ancova(2, 2, 0, a, b, tt) / c220^2
v_avg <- var_gamma_avg(2, 2, 0, a, b, tt) / c220^2

cat(sprintf('J_0=2, J_1=2 (sigma_b/sigma = 1):\n'))

```

```
154 ## J_0=2, J_1=2 (sigma_b/sigma = 1):
```

```
cat(sprintf(' GLS:    %.4f\n', v_gls))
```

```
155 ## GLS:    2.1569
```

```
cat(sprintf(' ANCOVA: %.4f\n', v_anc))
```

```
## ANCOVA: 2.9630
```

```
cat(sprintf(' Avg: %.4f\n', v_avg))
```

```
## Avg: 8.8889
```

```
cat(sprintf(' RE(GLS/ANC): %.3f\n', v_gls / v_anc))
```

```
## RE(GLS/ANC): 0.728
```

```
cat(sprintf(' RE(ANC/Avg): %.3f\n', v_anc / v_avg))
```

```
## RE(ANC/Avg): 0.333
```

2.4 SUR (seemingly unrelated regression)

The SUR approach treats the run-in and treatment periods as separate regression equations linked by correlated errors:

$$Y_1 = X_2 \begin{pmatrix} \beta_{11} \\ \beta_{12} \end{pmatrix} + e_1 \quad (\text{both groups}) \quad (7)$$

$$Y_2 = X_1 \alpha_2 + e_2 \quad (\text{baseline}) \quad (8)$$

The deviation-from-mean parameterization gives $\beta_1 = (\delta_1 + \delta_2)/2$ and $\beta_2 = (\delta_2 - \delta_1)/2$, where δ_j are the group-specific rates.

The SUR variance derivation and its comparison to the GLS and averaged estimators are not developed in this paper; the subsection is retained as a specification only and is flagged as future work in the Discussion.

2.5 BLUP as theoretical benchmark

The three estimators above are linear functions of the observed change scores and can be compared to a common benchmark: the best linear unbiased predictor (BLUP) of the individual random slopes. Under the random-slopes model with marginal covariance $V_i = R_i + Z_i G Z_i'$, the BLUP of the subject-specific random slope b_i given the data is

$$\hat{b}_i = G Z_i' V_i^{-1} (y_i - X_i \hat{\beta}) \quad (9)$$

where $\hat{\beta}$ is the GLS estimator of the fixed effects^{26,27}.

The group-specific slope estimators derived from the fixed effects share the same variance as $(\hat{\beta})_\gamma$, and the GLS estimator $\hat{\gamma}_{\text{GLS}}$ minimizes the variance among all linear unbiased estimators of γ when V_i is correctly specified. The averaged and ANCOVA estimators are specific linear approximations to this benchmark:

- **Averaged change-score** replaces V_i^{-1} with an unweighted sum; this is optimal only when $V_i \propto I$, which fails for any non-zero random-slope variance.
- **ANCOVA** collapses the run-in visits into a single scalar covariate (the run-in mean) rather than weighting them by V_i^{-1} ; the resulting efficiency loss vanishes as $J_0 \rightarrow \infty$ since the run-in mean becomes a sufficient statistic for the subject-specific intercept plus average slope.

Section 4 examines whether model misspecification erodes the nominal optimality of GLS relative to these approximations, which depend on fewer variance components. In the G and R misspecifications studied there, the GLS variance increases but remains below that of the averaged estimator; whether feasible GLS with estimated components loses its advantage in small samples is not studied here.

3 Relative efficiency under correct specification

3.1 Analytical comparison

The relative efficiency of the averaged estimator to GLS is:

$$\text{RE}_{\text{avg}} = \frac{\text{Var}(\hat{\gamma}_{\text{GLS}})}{\text{Var}(\hat{\gamma}_{\text{avg}}/c)} \leq 1 \quad (10)$$

with equality when the averaged estimator happens to coincide with GLS (which occurs in special cases).

The ordering asserted above can be stated and proved directly. One preliminary is needed. As defined, the averaged and ANCOVA contrasts average the accumulated treatment effect over the post-randomization visits, so their expectation is $c\gamma$ rather than γ , where

$$c = t \frac{J_1(J_1 + 1)/2 + J_1 J_2}{J_1 + J_2} \quad (11)$$

is the mean of the accumulated-effect design column over the post-randomization visits ($c = 1.5t$ for the $(2, 2, 0)$ design; $c = t$ when $J_1 = 1$, $J_2 = 0$). Dividing either contrast by c yields an unbiased estimator of γ with variance scaled by $1/c^2$.

Proposition. Under the random-slopes model with correctly specified Σ , $\text{Var}(\hat{\gamma}_{\text{GLS}}) \leq \text{Var}(\hat{\gamma}_{\text{ANC}}/c) \leq \text{Var}(\hat{\gamma}_{\text{avg}}/c)$.

Proof. Each rescaled contrast is a linear unbiased estimator of γ , and $\hat{\gamma}_{\text{GLS}}$ is the best linear unbiased estimator of γ when Σ is known (Gauss-Markov), which gives the first inequality. For the second, both contrasts have the form (post-randomization mean) $-\phi \times$ (run-in mean): the population ANCOVA coefficient minimizes the residual variance over ϕ , while the averaged estimator fixes $\phi = 1$, so the ANCOVA contrast variance is no larger; dividing both by c^2 preserves the inequality. (The feasible ANCOVA estimates ϕ and carries the usual finite-sample inflation.) $\square$

The ordering itself is not new.⁵ established the same qualitative ranking for repeated-measures

Table 1: RE of averaged change-score vs GLS on the γ scale ($\text{Var}_{\text{GLS}} / \text{Var}(\hat{\gamma}_{\text{avg}}/c)$) by σ_b/σ .

(J_0, J_1, J_2)	0.5	1	2	5	20
(1,1,0)	0.667	0.500	0.250	0.056	0.004
(2,1,0)	0.653	0.344	0.117	0.021	0.001
(4,1,0)	0.427	0.149	0.041	0.007	0.000
(0,2,0)	0.818	0.900	0.964	0.994	1.000
(1,2,0)	0.529	0.410	0.204	0.045	0.003
(2,2,0)	0.490	0.243	0.080	0.014	0.001
(4,2,0)	0.276	0.089	0.024	0.004	0.000
(0,3,0)	0.771	0.900	0.969	0.995	1.000
(1,3,0)	0.501	0.394	0.193	0.042	0.003
(2,3,0)	0.433	0.209	0.068	0.012	0.001
(4,3,0)	0.220	0.069	0.018	0.003	0.000
(0,1,1)	0.727	0.533	0.245	0.050	0.003
(1,1,1)	0.452	0.258	0.093	0.017	0.001
(2,1,1)	0.422	0.181	0.055	0.009	0.001
(4,1,1)	0.255	0.080	0.021	0.003	0.000
(0,2,1)	0.705	0.517	0.224	0.044	0.003
(1,2,1)	0.448	0.262	0.094	0.017	0.001
(2,2,1)	0.401	0.172	0.052	0.009	0.001
(4,2,1)	0.221	0.068	0.018	0.003	0.000
(0,3,1)	0.685	0.488	0.202	0.039	0.003
(1,3,1)	0.446	0.258	0.092	0.017	0.001
(2,3,1)	0.380	0.162	0.049	0.008	0.001
(4,3,1)	0.194	0.059	0.016	0.003	0.000

designs with several baseline measurements, and⁶ extended it to repeated post-treatment measurements; the argument above is the specialization of that result to the run-in random-slope model, where the run-in mean plays the role of their baseline mean. The contribution of the following sections is the closed-form magnitude of the two gaps as functions of (J_0, J_1, J_2) and σ_b/σ , which the Frison-Pocock parameterization does not deliver for a slope estimand or for designs with a common close phase.

Numerically, the ordering holds with zero violations over the 115-cell grid of the tables below. All tables, figures, and simulation results in this paper are on the equal-estimand γ scale: the averaged and ANCOVA contrast variances are divided by c^2 before any comparison. For reference, the raw phase-mean contrasts carry variances larger by the factor c^2 ($c^2 = 2.25$ at the $(2, 2, 0)$ design), so raw-contrast comparisons would overstate the penalty of the simpler estimators.

3.2 When does the gap exceed 5%?

The answer, computed over the full grid of the preceding tables (all on the γ scale), is: in most cells. Of the 115 design-by-ratio cells, 12 have ANCOVA relative efficiency of at least 0.95

Table 2: RE of ANCOVA vs GLS on the γ scale ($\text{Var}_{\text{GLS}} / \text{Var}(\hat{\gamma}_{\text{ANC}}/c)$) by σ_b/σ .

(J_0, J_1, J_2)	0.5	1	2	5	20
(1,1,0)	1.000	1.000	1.000	1.000	1.000
(2,1,0)	0.971	0.909	0.860	0.838	0.834
(4,1,0)	0.838	0.747	0.713	0.702	0.700
(0,2,0)	0.818	0.900	0.964	0.994	1.000
(1,2,0)	0.858	0.866	0.853	0.838	0.834
(2,2,0)	0.808	0.728	0.671	0.648	0.643
(4,2,0)	0.637	0.544	0.512	0.502	0.500
(0,3,0)	0.771	0.900	0.969	0.995	1.000
(1,3,0)	0.808	0.815	0.785	0.758	0.751
(2,3,0)	0.732	0.642	0.582	0.558	0.553
(4,3,0)	0.538	0.450	0.420	0.411	0.409
(0,1,1)	0.727	0.533	0.245	0.050	0.003
(1,1,1)	0.733	0.545	0.390	0.319	0.304
(2,1,1)	0.697	0.542	0.460	0.430	0.425
(4,1,1)	0.587	0.490	0.456	0.446	0.445
(0,2,1)	0.705	0.517	0.224	0.044	0.003
(1,2,1)	0.723	0.541	0.385	0.313	0.299
(2,2,1)	0.678	0.527	0.447	0.418	0.412
(4,2,1)	0.539	0.444	0.413	0.403	0.401
(0,3,1)	0.685	0.488	0.202	0.039	0.003
(1,3,1)	0.699	0.516	0.362	0.293	0.279
(2,3,1)	0.642	0.494	0.416	0.388	0.382
(4,3,1)	0.486	0.397	0.367	0.359	0.357

(an efficiency loss below 5%), and they form exactly two families. The first (6 cells) comprises single-post-visit run-in designs with no common close ($J_1 = 1, J_2 = 0$): the (1, 1, 0) design at every tabulated ratio, plus (2, 1, 0) at $\sigma_b/\sigma = 0.5$; with a single post-randomization visit there is no within-period slope information for GLS to exploit. The second (6 cells) comprises designs without run-in or common close ($J_0 = 0, J_2 = 0, J_1 \in \{2, 3\}$) at $\sigma_b/\sigma \geq 2$, where ANCOVA reduces to the averaged estimator and the rescaled phase mean is nearly efficient. In run-in designs with $J_1 \geq 2$, the regime that motivates this paper, the loss is substantial: at the representative (2, 2, 0) design the relative efficiency ranges from 0.81 at $\sigma_b/\sigma = 0.5$ down to 0.64 at $\sigma_b/\sigma = 20$, an efficiency loss of 19% to 36%. The averaged estimator reaches 0.95 only in the second family (6 of 115 cells).

4 Robustness to misspecification

4.1 Misspecified G

When σ_b^2 is misestimated (e.g., too large or too small), the GLS analysis that assumes it is no longer optimal. Its sampling variance is then the sandwich form $(X'\Sigma_a^{-1}X)^{-1}X'\Sigma_a^{-1}\Sigma\Sigma_a^{-1}X(X'\Sigma_a^{-1}X)^{-1}$, with Σ_a the assumed and Σ the true covariance, which is the model-robust variance of²⁸ evaluated at a known working covariance rather than an estimated one. The averaged change-score estimator's definition does not depend on the assumed value, so its variance is constant across the sweep. Note the scale of the comparison: over assumed $\sigma_b^2 \in [0.01, 5]$ with true $\sigma_b^2 = 1$, the misspecified GLS sandwich variance never exceeds 5.30, still below the averaged-estimator variance of 8.89 on the γ scale. The studied misspecification erodes but does not overturn the GLS advantage.

4.2 Misspecified R

When the residual correlation structure is wrong (e.g., compound symmetry assumed but AR(1) is true), both the GLS variance formula and the Woodbury correction are biased. The averaged estimator's definition does not depend on the assumed model, though its γ -scale variance does change with the true R (8.4 to 9.1 across $\rho \in [0, 0.9]$ in the figure). The GLS sandwich variance ranges from 0.74 to 2.16 over the same sweep, so the studied R misspecification leaves GLS four to eleven times more efficient than the averaged estimator throughout.

5 Simulation study

5.1 Design

To complement the analytical results, we simulate trials under the following factorial design:

- $m \in \{50, 100\}$ per group
- $(J_0, J_1, J_2) \in \{(0, 2, 0), (2, 2, 0), (2, 2, 2)\}$

Misspecified random slope variance




Figure 1: (#fig:misspec_G)Effect of misspecifying σ_b^2 on GLS variance. True $\sigma_b^2 = 1$. The GLS estimator uses an assumed value on the x-axis. The averaged estimator's definition does not use σ_b^2 , so its γ -scale variance (horizontal line) does not depend on the assumed value. The misspecified GLS variance remains below the averaged variance throughout.

CS assumed, AR(1) true




Figure 2: (#fig:misspec_R)Effect of analyzing under compound symmetry when the true correlation is AR(1). Sandwich variance of GLS (solid) vs averaged change-score on the γ scale (dashed) as a function of ρ . The averaged estimator's definition does not use the assumed correlation, but its variance varies with the true R ; GLS remains at least fourfold more efficient at every ρ .

Table 3: Monte Carlo comparison of GLS (correctly specified), GLS (misspecified σ_b^2 , assuming $\sigma_b^2 = \sigma^2$), and the averaged change-score estimator under compound symmetry (1000 replicates, $\Delta = 0.5$, two-sided $\alpha = 0.05$). Columns report the empirical standard deviation of $\hat{\gamma}$ and the empirical power. The averaged estimator is rescaled to the γ scale (contrast divided by c).

Design	σ_b/σ	m	SD GLS	SD mis	SD avg	Pow GLS	Pow mis	Pow avg
(0,2,0)	1	50	0.238	0.238	0.250	0.537	0.537	0.495
(2,2,0)	1	50	0.210	0.210	0.417	0.677	0.677	0.206
(2,2,2)	1	50	0.151	0.151	0.458	0.902	0.902	0.192
(0,2,0)	1	100	0.169	0.169	0.178	0.823	0.823	0.769
(2,2,0)	1	100	0.146	0.146	0.308	0.926	0.926	0.393
(2,2,2)	1	100	0.108	0.108	0.334	0.997	0.997	0.345
(0,2,0)	5	50	1.012	1.012	1.017	0.083	0.653	0.083
(2,2,0)	5	50	0.241	0.528	2.043	0.555	0.600	0.067
(2,2,2)	5	50	0.164	0.284	2.297	0.873	0.765	0.054
(0,2,0)	5	100	0.726	0.726	0.726	0.108	0.733	0.109
(2,2,0)	5	100	0.162	0.360	1.387	0.857	0.748	0.067
(2,2,2)	5	100	0.106	0.189	1.573	0.997	0.936	0.060

- $\sigma_b/\sigma \in \{1, 5\}$
- True correlation: compound symmetry (an AR(1) truth is examined analytically in Section 4 but is not simulated)
- Analysis: GLS (correctly specified), GLS (misspecified), averaged change-score

For each configuration, 1000 trials are simulated. The outcomes are the empirical standard deviation of $\hat{\gamma}$ and the empirical power. Coverage of nominal 95% confidence intervals and empirical size at $\Delta = 0$ are not assessed; extending the factorial to $m \in \{30, 200\}$, $\sigma_b/\sigma = 20$, an AR(1) truth arm, and 5000 replicates remains future work. Test statistics use standard normal critical values rather than t ; the distinction is immaterial at $m \geq 50$ per group.

5.2 Results

We simulate the factorial design described above under compound symmetry, comparing the correctly specified GLS estimator, a GLS estimator that misstates the random-slope variance (it assumes $\sigma_b^2 = \sigma^2$ regardless of the truth), and the averaged change-score estimator, the latter rescaled to the γ scale by dividing the contrast by c . For each cell we report the empirical standard deviation of $\hat{\gamma}$ and the empirical power from 1000 replicates; the Monte Carlo standard error of a tabulated power estimate is at most $\sqrt{0.25/1000} \approx 0.016$.

One feature of the table requires comment. In the (0, 2, 0) rows the correctly specified and misspecified GLS analyses report identical empirical standard deviations but different powers. With $J_0 = 0$ the design has two observations and two fixed effects, so the GLS point estimator

Table 4: Sample size per group: GLS, ANCOVA, and averaged estimators (MIRIAD parameters, $\Delta = 0.25$, 90% power). All variances are on the γ scale (ANCOVA and averaged contrasts divided by c).

(J_0, J_1, J_2)	N GLS	N ANC	N Avg
(4,3,1)	36	95	1630
(4,3,0)	43	100	1676
(4,2,1)	60	144	2413
(2,3,1)	69	162	1053
(4,2,0)	78	151	2368
(2,3,0)	93	157	1028
(1,3,1)	99	255	821
(2,2,1)	104	226	1479
(0,3,1)	112	433	433
(1,2,1)	138	335	1116
(2,2,0)	145	213	1361
(0,2,1)	150	527	527
(4,1,1)	156	339	5328
(1,3,0)	185	233	775
(2,1,1)	226	478	3063

does not depend on the assumed covariance and the two analyses coincide as estimators; the apparent power advantage of the misspecified analysis reflects only an anticonservative model-based standard error, not additional information. More generally, because empirical size at $\Delta = 0$ and interval coverage are not tabulated, power comparisons across variance models in this table are not interpretable as efficiency comparisons; adding size and coverage columns is remaining work.

6 Numerical examples

6.1 MIRIAD parameters

6.2 When to use which estimator

A decision rule based on m , σ_b/σ , and confidence in the correlation structure requires the feasible-GLS simulation described in the Discussion and is deferred to a revision. The results above support only the narrower statement that, with variance components taken as known, GLS dominates and the penalty of the simpler estimators is largest in run-in designs with $J_1 \geq 2$ and high σ_b/σ .

7 Discussion

The estimator comparison yields a graded recommendation rather than a single winner. Under correct specification GLS is efficient and dominates both ANCOVA and the averaged change-score estimator, as the ordering $\text{Var}_{\text{GLS}} \leq \text{Var}_{\text{ANC}} \leq$

306 Var_{avg} confirms. This dominance is a large-sample statement: it presumes the
307 covariance components are known or well estimated. With moderate-to-large sam-
308 ples and a credible covariance model, GLS is the appropriate default.

309 The misspecification analysis qualifies the case for GLS less than might be ex-
310 pected. Because GLS weights the data by an assumed inverse covariance, errors
311 in the assumed G or R feed into its sampling variance, and the sandwich-variance
312 calculations quantify the inflation. In the settings studied, however, the inflation
313 never closes the gap: the misspecified GLS variance peaks at 5.30 in the G sweep
314 and at 2.16 in the R sweep, against γ -scale averaged-estimator variances of 8.4 to
315 9.1 at the same design. The averaged change-score estimator avoids the variance
316 components entirely, and ANCOVA estimates the run-in coefficient from the data,
317 but under the misspecifications examined neither robustness property is enough to
318 overcome their efficiency deficit. These calculations take the variance components
319 of the GLS analysis as fixed inputs; the open and practically important question is
320 whether feasible GLS, with components estimated by REML from small samples,
321 keeps its advantage. That question is not answered here.

322 The practical importance of the estimator choice depends on the design. As the
323 numbers of run-in and treatment observations grow, the phase means approach
324 sufficiency for the subject-specific slope, the ANCOVA efficiency loss vanishes as
325 $J_0 \rightarrow \infty$, and the relative efficiencies of all three estimators converge toward one.
326 The gap is largest, and the choice of estimator therefore most consequential, in
327 small designs with few observations per phase and a high slope-to-noise ratio,
328 exactly the regime in which the estimated covariance is also least stable. Whether
329 that instability erodes the feasible-GLS advantage enough to change the ranking
330 in small studies is the question the planned feasible-GLS simulation must answer;
331 the known-components results here cannot.

332 Two estimators not fully developed here would round out the comparison. The
333 seemingly unrelated regression formulation of Section 2.4 treats the run-in and
334 treatment periods as linked equations and offers a compromise between the aver-
335 aged estimator and full GLS, at the cost of additional implementation complexity;
336 its variance derivation remains to be completed. The constrained longitudinal
337 data analysis estimator uses the full covariance without committing to a random-
338 effects distribution and should fall between ANCOVA and GLS in efficiency. A
339 comprehensive treatment would compare all of these under a feasible-GLS simula-
340 tion in which the variance components are estimated rather than assumed, which
341 is the setting in which the robustness distinctions drawn here become decisive.

342 8 Conflict of interest

343 The author declares no conflicts of interest.

344 9 Funding

345 This work received no specific funding.

346 10 Data availability statement

347 The `runinpower` R package implementing the GLS, ANCOVA, and averaged-
348 change-score estimators (`ancova.R`, `averaged.R`), the simulation drivers, and the
349 shared bibliography across the four-paper compendium are openly available in the
350 project repository. Specific paths are listed in the Reproducibility section below.

351 11 Reproducibility

352 This manuscript was rendered with R 4.4.x. Principal packages used are the
353 in-package `runinpower` source, `tinytest` (testing harness), and `rmarkdown` plus
354 `bookdown` for the manuscript build. The GLS comparator in the simulation is
355 implemented directly in the manuscript source (`analyze_gls` in the simulation
356 chunk); no external mixed model fitter is used.

357 **Random-number generator.** The Monte Carlo simulation comparing the three
358 estimators sets `set.seed(20260418)` once, immediately before the replicate loop.
359 No per-replicate reseeding is performed, so the full replicate stream, and hence
360 the tabulated results, follow deterministically from that single seed.

361 **Data and code availability.** The R package source is at `R/`. Simulation
362 drivers are at `analysis/scripts/`. The companion manuscripts (Papers 1, 2,
363 3) are at `analysis/report/{01,02,03}-*/`. The manuscript source and shared
364 bibliography (`references.bib`, a symlink to `../../references.bib`) are at
365 `analysis/report/04-gls-vs-ttest/`.

366 12 References

- 367 1. Senn S. Change from baseline and analysis of covariance revisited. *Statistics in Medicine*. 2006;25(24):4334-4344. doi:10.1002/sim.2682
- 368 2. Van Breukelen GJP. ANCOVA versus change from baseline had more power in randomized studies and more bias in nonrandomized studies. *Journal of Clinical Epidemiology*. 2006;59(9):920-925. doi:10.1016/j.jclinepi.2006.02.007
- 369 3. Crager MR. Analysis of covariance in parallel-group clinical trials with pretreatment baselines. *Biometrics*. 1987;43(4):895-901. doi:10.2307/2531542

- 370 4. Yang L, Tsiatis AA. Efficiency study of estimators for a treatment effect in a pretest-posttest trial. *The American Statistician*. 2001;55(4):314-321. doi:10.1198/000313001753272466
- 371 5. Frison L, Pocock SJ. Repeated measures in clinical trials: Analysis using mean summary statistics and its implications for design. *Statistics in Medicine*. 1992;11(13):1685-1704. doi:10.1002/sim.4780111304
- 372 6. Frison L, Pocock SJ. Linearly divergent treatment effects in clinical trials with repeated measures: Efficient analysis using summary statistics. *Statistics in Medicine*. 1997;16(24):2855-2872. doi:10.1002/(SICI)1097-0258(19971230)16:24<2855::AID-SIM749>3.0.CO;2-2
- 373 7. Hu N, Mackey H, Thomas R. Power and sample size for random coefficient regression models in randomized experiments with monotone missing data. *Biometrical Journal*. 2021;63(4):806-824. doi:10.1002/bimj.202000184
- 374 8. Liu GF, Lu K, Mogg R, Mallick M, Mehrotra DV. Should baseline be a covariate or dependent variable in analyses of change from baseline in clinical trials? *Statistics in Medicine*. 2009;28(20):2509-2530. doi:10.1002/sim.3639
- 375 9. Lu K. On efficiency of constrained longitudinal data analysis versus longitudinal analysis of covariance. *Biometrics*. 2010;66(3):891-896. doi:10.1111/j.1541-0420.2009.01322.x
- 376 10. Tsiatis AA, Davidian M, Zhang M, Lu X. Covariate adjustment for two-sample treatment comparisons in randomized clinical trials: A principled yet flexible approach. *Statistics in Medicine*. 2008;27(23):4658-4677. doi:10.1002/sim.3113
- 377 11. Mallinckrodt CH, Ye J, et al. Mixed models for repeated measures should include time-by-covariate interactions to assure power gains and robustness against dropout bias relative to complete-case ANCOVA. *Statistics in Biopharmaceutical Research*. Published online 2021. doi:10.1080/19466315.2021.1976846
- 378 12. Brogan DR, Kutner MH. Comparative analyses of pretest-posttest research designs. *The American Statistician*. 1980;34(4):229-232. doi:10.1080/00031305.1980.10483034
- 379 13. Coffman CJ, Edelman D, Woolson RF. To condition or not condition? Analysing “change” in longitudinal randomised controlled trials. *BMJ Open*. 2016;6(12):e013096. doi:10.1136/bmjopen-2016-013096

- 380 14. Wang G, Aschenbrenner AJ, Li Y, et al. Two-period linear mixed effects models to analyze clinical trials with run-in data when the primary outcome is continuous: Applications to Alzheimer's disease. *Alzheimer's & Dementia: Translational Research & Clinical Interventions*. 2019;5:450-457. doi:10.1016/j.trci.2019.08.002
- 381 15. Liang KY, Zeger SL. Longitudinal data analysis of continuous and discrete responses for pre-post designs. *Sankhya: The Indian Journal of Statistics, Series B*. 2000;62(1):134-148.
- 382 16. Funatogawa I, Funatogawa T. Longitudinal analysis of pre- and post-treatment measurements with equal baseline assumptions in randomized trials. *Biometrical Journal*. 2020;62(5):1235-1249. doi:10.1002/bimj.201800389
- 383 17. Iddi S, Donohue MC. Power and sample size for longitudinal models in R – the long-power package and Shiny app. *The R Journal*. 2022;14(1):264-282. <https://journal.r-project.org/articles/RJ-2022-022/>
- 384 18. Liu G, Liang KY. Sample size calculations for studies with correlated observations. *Biometrics*. 1997;53:937-947. doi:10.2307/2533554
- 385 19. Diggle PJ, Heagerty P, Liang KY, Zeger SL. *Analysis of Longitudinal Data*. 2nd ed. Oxford University Press; 2002.
- 386 20. Ard MC, Edland SD. Power calculations for clinical trials in Alzheimer's disease. *Journal of Alzheimer's Disease*. 2011;21:369-377. doi:10.3233/JAD-2011-0062
- 387 21. Frost C, Kenward MG, Fox NC. Optimizing the design of clinical trials where the outcome is a rate. Can estimating a baseline rate in a run-in period increase efficiency? *Statistics in Medicine*. 2008;27(18):3717-3731. doi:10.1002/sim.3280
- 388 22. Laird NM, Ware JH. Random-effects models for longitudinal data. *Biometrics*. 1982;38(4):963-974. doi:10.2307/2529876
- 389 23. Verbeke G, Molenberghs G. *Linear Mixed Models for Longitudinal Data*. Springer; 2000.
- 390 24. Borm GF, Fransen J, Lemmens WAJG. A simple sample size formula for analysis of covariance in randomized clinical trials. *Journal of Clinical Epidemiology*. 2007;60(12):1234-1238. doi:10.1016/j.jclinepi.2007.02.006
- 391 25. Vickers AJ. The use of percentage change from baseline as an outcome in a controlled trial is statistically inefficient: A simulation study. *BMC Medical Research Methodology*. 2001;1:6. doi:10.1186/1471-2288-1-6

- 392 26. Henderson CR. Best linear unbiased estimation and prediction under a selection model. *Biometrics*. 1975;31(2):423-447. doi:10.2307/2529430
- 393 27. Robinson GK. That BLUP is a good thing: The estimation of random effects. *Statistical Science*. 1991;6(1):15-32. doi:10.1214/ss/1177011926
- 394 28. Liang KY, Zeger SL. Longitudinal data analysis using generalized linear models. *Biometrika*. 1986;73(1):13-22. doi:10.1093/biomet/73.1.13